

# IoA Project - MNIST\_R

*Aravind*

*March 13, 2018*

## DIGIT RECOGNIZER - IOA PROJECT

Read data : <https://www.kaggle.com/c/digit-recognizer>

```
train <- read.csv("C:/Users/Aravind/Documents/Digit_Recognizer/train.csv")
test  <- read.csv("C:/Users/Aravind/Documents/Digit_Recognizer/test.csv")
```

## Dimensions of train and test

```
dim(train)
```

```
## [1] 42000  785
```

```
dim(test)
```


```
## [1] 28000  784
```

```
train$label <- as.factor(train$label)
```

## Visualization


1st Image

```
img<-matrix((train[1,2:ncol(train)]), nrow=28, ncol=28) #For the 1st Image
img_numbers <- apply(img, 2, as.numeric)
image(1:28, 1:28, img_numbers, col=gray((0:255)/255))
```




2nd Image

```
img<-matrix((train[2,2:ncol(train)]), nrow=28, ncol=28) #For the 2nd Image  
img_numbers <- apply(img, 2, as.numeric)  
image(1:28, 1:28, img_numbers, col=gray((0:255)/255))
```




For the first ten rows

```
par(mfrow=c(4, 3))
for (i in 1:10){
  img<-matrix((train[i,2:ncol(train)]), nrow=28, ncol=28)
  img_numbers <- apply(img, 2, as.numeric)
  image(1:28, 1:28, img_numbers, col=gray((0:255)/255))
}
```



Average image of each digit

```
par(mfrow=c(4,3))
img<-array(dim=c(10,28*28))
for(i in 0:9){
  img[i+1,<-apply(train[train[,1]==i,-1],2,sum)
  img[i+1,<-img[i+1,]/255*255
  im<-array(img[i+1,],dim=c(28,28))
  im<-im[,28:1] #right side up
  image(1:28,1:28,im,col = grey(0:255/255),main=i)
}
```



## Split train data

```
require(caret)

## Loading required package: caret
## Loading required package: lattice
## Loading required package: ggplot2

set.seed(123)
index <- createDataPartition(train$label, p=0.80, list = F)
train_set <- train[index,]
```

```
test_set <- train[-index,]
```

## Scaling and Centering


```
X <- train_set[, -1]
X_scale <- X/255
X_center <- scale(X_scale, center = T, scale = F)
Y <- test_set[, -1]
Y_scale <- Y/255
Y_center <- scale(Y_scale, center = T, scale = F)
```

## Principal Component Analysis

```
pca<-princomp(X_center)
std_dev <- pca[1:250]$sdev
pr_var <- std_dev^2
prop_varex <- pr_var/sum(pr_var)
```

## Plot

```
plot(cumsum(prop_varex[1:250]), xlab = "Principal Component",
     ylab = "Cumulative Proportion of Variance Explained",
     type = "b")
```



Using first 250 components we can explain ~100% of variation

## Splitting PCA components into train and test

```
train_set_pca <- data.frame(predict(pca, newdata = train_set[, -1]))[1:250]
train_set_pca$label <- train_set$label
test_set_pca <- data.frame(predict(pca, newdata = test_set[, -1]))[1:250]
test_set_pca$label <- test_set$label
```

## Model Building

### DECISION TREE

```
require(rpart)
```

```
## Loading required package: rpart
```

```
model_rpart <- rpart(label ~., data = train_set_pca)
```

### Predict test\_\_set

```
pred_rpart <- predict(model_rpart, newdata = test_set_pca[-251], type = 'class')
```

### Confusion matrix

```
cm = table('Actual Digit' = test_set_pca[, 251], 'Predicted Digit' = pred_rpart)
cm
```

```
##          Predicted Digit
## Actual Digit  0  1  2  3  4  5  6  7  8  9
##          0 475  0  95 136  1  75 12 12  7 13
##          1  0 778  43  30  0  38 23  0 24  0
##          2  32  5 597  75  7  10 36  3 43 27
##          3  20  4  61 614  1  35 38  0 87 10
##          4  0 18  35  19 468 13 11 12 31 207
##          5  27  2 102 237 21 208 22 10 94  36
##          6  43  2 108  61  4  9 443 26 27 104
##          7  0 49  32  24 15  31  1 487 35 206
##          8  25  6  83  81  7  67 14  1 483  45
##          9  4 47  15  22 98  33  2  54 32 530
```

### Accuracy

```
print(sum(diag(cm))/sum(cm))
```

```
## [1] 0.6054073
```

Decision tree gives an accuracy of 60%

## Random Forest Model

```
require(randomForest)
```



```
## Loading required package: randomForest
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
model <- randomForest(label ~., data = train_set_pca, ntree = 100)
```

## Predict test\_set

```
pred <- predict(model, newdata = test_set_pca[-251])
```

## Confusion Matrix

```
cm = table('Actual Digit' = test_set_pca[, 251], 'Predicted Digit' = pred)
cm
```

```
##               Predicted Digit
## Actual Digit  0  1  2  3  4  5  6  7  8  9
##               0 807  0  2  1  0  4  7  0  5  0
##               1  0 914  5  2  1  4  5  2  2  1
##               2  6  3 777 12  8  3  2  3 18  3
##               3  3  0 16 792  2 23  5  5 16  8
##               4  3  3  1  0 772  3  3  2  5 22
##               5  3  1  3 28  6 697  8  2  5  6
##               6  4  0  5  0  3  3 810  0  2  0
##               7  1  2  3  0 10  2  1 830  8 23
##               8  1  8 10 21 10 16  9  1 728  8
##               9  6  2  2 11 29  7  3 16  4 757
```

## Accuracy

```
print(sum(diag(cm))/sum(cm))
```

```
## [1] 0.9390186
```

Random Forest gives an accuracy of 93.9%

## K-NN

```
require(class)

## Loading required package: class
model_knn_pred <- knn(train = train_set_pca[, -251],
                      test = test_set_pca[, -251],
                      cl = train_set_pca[, 251],
                      k=3)
```

## Confusion Matrix

```
cm_knn <- table('Actual Digit' = test_set_pca[, 251], 'Predicted Digit' = model_knn_pred)
cm_knn
```

```
##          Predicted Digit
## Actual Digit  0  1  2  3  4  5  6  7  8  9
##          0 819  1  0  0  0  2  3  0  1  0
##          1  0 929  2  0  0  0  1  3  0  1
##          2  6 10 796  1  1  0  1 17  2  1
##          3  1  3  7 838  0  7  1  2  8  3
##          4  0  7  0  0 783  0  1  0  0 23
##          5  0  0  1 16  1 726  9  0  3  3
##          6  4  0  0  0  1  1 821  0  0  0
##          7  0  7  1  0  1  0  0 858  0 13
##          8  3  9  6  6  5 10  3  1 754 15
##          9  1  0  2  4 14  1  1  9  1 804
```

## Accuracy

```
print(sum(diag(cm_knn))/sum(cm_knn))
```

```
## [1] 0.96808
```

K-NN gives an accuracy of 96.8%

---

## SVM

```
require(e1071)

## Loading required package: e1071
model_svm = svm(formula = label ~ .,
                 data = train_set_pca,
                 type = 'C-classification',
                 kernel = 'linear')
```

```
pred_svm <- predict(model_svm, newdata = test_set_pca[-251])
```

## Confusion Matrix

```
cm_svm = table('Actual Digit' = test_set_pca[, 251], 'Predicted Digit' = pred_svm)
cm_svm
```

```
##          Predicted Digit
## Actual Digit  0  1  2  3  4  5  6  7  8  9
##          0 797  0  4  0  3  6  9  0  7  0
##          1  0 920  4  3  0  2  1  2  4  0
##          2 14 13 722 17 18  7 10 11 21  2
##          3  5  7 14 783  3 26  0 13 15  4
##          4  0  3  5  0 771  1  7  5  4 18
##          5  9  4  6 29  7 675  9  3 13  4
##          6  7  3  7  4  4 13 788  1  0  0
##          7  2  4  6  6  9  1  0 825  3 24
##          8  7 19  5 21  7 24  1  1 723  4
##          9  4  2  3  4 39  6  0 31  9 739
```

## Accuracy

```
print(sum(diag(cm_svm))/sum(cm_svm))
```

```
## [1] 0.9222249
```

SVM gives an accuracy of 92%

---

## XG BOOST

```
require(xgboost)
```

```
## Loading required package: xgboost
```

```
model_xgb <- xgboost(data = as.matrix(train_set_pca[-251]), label = train_set_pca$label, nrounds = 500)
```

By using nround = 500, we are able to reduce rmse value from 4.2 to 0.16. The lesser it is the better your model performs.

## Prediction

```
pred_xgb <- predict(model_xgb, newdata = as.matrix(test_set_pca[-251]))
pred_xgb <- (pred_xgb >= 0.5)
```

## Confusion Matrix

```
cm_xgb <- table('Actual Digit' = test_set_pca[, 251], 'Predicted Digit' = pred_xgb)
cm_xgb
```

```
##           Predicted Digit
## Actual Digit FALSE TRUE
##           0   126  700
##           1     1  935
##           2     2  833
##           3     0  870
##           4     0  814
##           5     0  759
##           6     1  826
##           7     0  880
##           8     0  812
##           9     0  837
```

## Accuracy

```
print(sum(cm_xgb[,2])/sum(cm_xgb))
```

```
## [1] 0.9845164
```

XGBOOST gives an accuracy of 98.45%

## Plot accuracy of each digits

```
total <- apply(cm_xgb,1, sum)
total <- array(total)
total
```

```
## [1] 826 936 835 870 814 759 827 880 812 837
```

```
crt <- c()
for (i in 1:10){
  crt[i] <- cm_xgb[i,2]
}
crt
```

```
## [1] 700 935 833 870 814 759 826 880 812 837
```

```
result <- crt/total
result
```

```
## [1] 0.8474576 0.9989316 0.9976048 1.0000000 1.0000000 1.0000000 0.9987908
## [8] 1.0000000 1.0000000 1.0000000
```

```
digits <- c(0:9)
digits
```

```
## [1] 0 1 2 3 4 5 6 7 8 9
```

```

result_df <- data.frame(digits = digits, accuracy = result)
result_df$digits <- as.factor(result_df$digits)
result_df$accuracy <- result_df$accuracy*100
result_df$accuracy <- round(result_df$accuracy, digits = 2)
result_df

```

```


##      digits accuracy
## 1         0    84.75
## 2         1    99.89
## 3         2    99.76
## 4         3   100.00
## 5         4   100.00
## 6         5   100.00
## 7         6    99.88
## 8         7   100.00
## 9         8   100.00
## 10        9   100.00

```

```

require(ggplot2)
ggplot(result_df, aes(digits,accuracy, fill = digits)) + geom_bar(colour="black",stat = "identity") + c

```



The barplot shows the model struggles to identify the digit 0.

## Lollipop Chart

```
require(ggplot2)
ggplot(result_df, aes(x=digits, y=accuracy, fill = digits)) +
  geom_segment(aes(x=digits, xend=digits, y=80, yend=accuracy), size=2, color="blue", linetype="solid") +
  geom_point(size=10, color="pink", fill=alpha("red", 0.3), alpha=0.7, shape=21, stroke=2) +
  theme_light()
```

